

The Dr. X Files: Secrets, Patterns, and Hidden Insights

Date: 17th April 2025

Submitted by: Said Al-Mujaini

1. The Mission & Ground Rules

Welcome to the Dr. X Analysis Suite! Our goal is to dissect Dr. X's documents (.pdf, .docx, .xlsx, etc.) using pure Natural Language Processing. We're digging for insights, maybe even clues to their disappearance.

Crucially: We operate entirely locally – local LLMs (like Llama 2 via Ollama), local vector storage (ChromaDB).

1.1. Tech Toolkit Highlights

- **Orchestrator:** Langchain
- **Brain:** Local LLMs (via langchain-ollama)
- **Understanding:** nomic-embed-text-v1 Embeddings (sentence-transformers)
- **Memory:** chromadb (Persistent Vector Store)
- **Parsing Power:** PyMuPDF, python-docx, pandas
- **Metrics:** rouge-score, Custom Performance Logging

1.2. The Big Picture: System Flow

Everything starts with Dr. X's files and flows through modular processing stages, tailored for different tasks like Q&A, Translation, or Summarization. Configuration (config.py) guides the process.

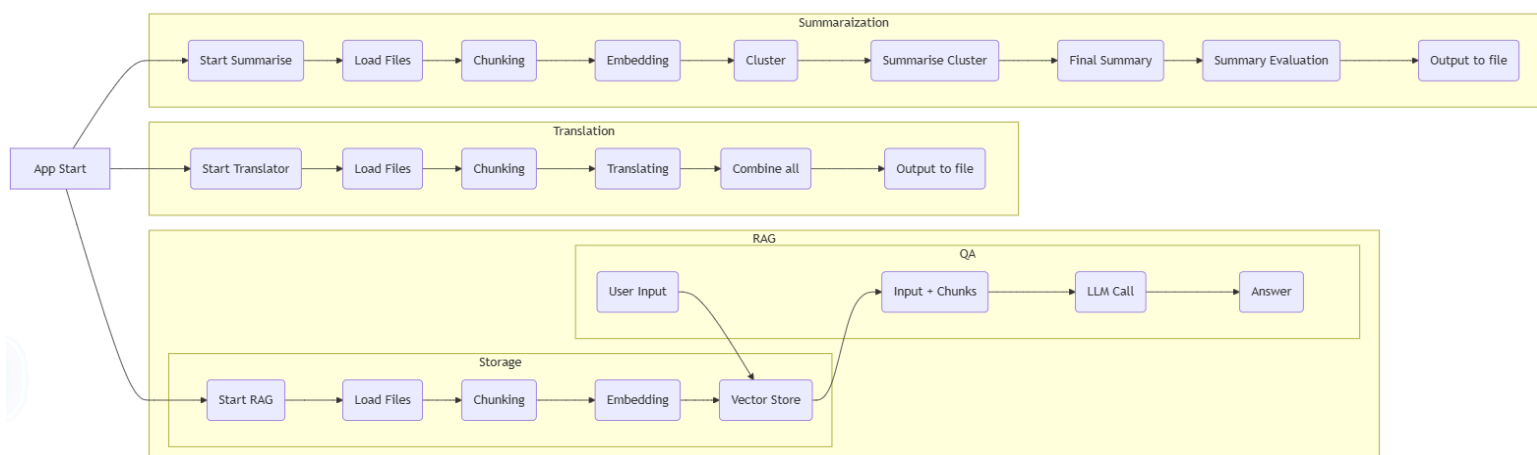


Figure 1: High-Level System Architecture

1.3. Central Control: config.py

This file is mission control! It defines model names, file paths, chunk sizes, vector DB settings, and crucial prompt templates, allowing easy tuning without touching the core code.

2. Decoding the Documents

Before analysis, we need clean, structured text. This involves smart loading and strategic chunking.

2.1. Smart Loading (EnhancedDocumentLoader)

We don't just grab text; we understand structure, especially tables.

- **Workflow:**
 1. **Detecting Format:** Identify .pdf, .docx, .xlsx, etc.
 2. **Extract Content:** Use specialized libraries (PyMuPDF, python-docx, pandas).
 3. **Spot Tables:** Find tables within PDFs, DOCX and XLSX files.
 4. **Translate Tables:** This is key! Use `_format_table` to convert raw table data (lists of lists) into human-readable sentences. *Why?* LLMs understand "The record shows Capacity is 100 units" better than just "100".
 5. **Combine & Package:** Merge text and formatted table descriptions into Langchain Document objects with rich metadata (source, page/row).

2.2. Strategic Chunking (Chunker)

Large documents overwhelm LLMs. We break them down intelligently using the **cl100k_base** tokenizer.

- **Workflow:** Takes Document objects -> Outputs smaller text chunks with metadata.

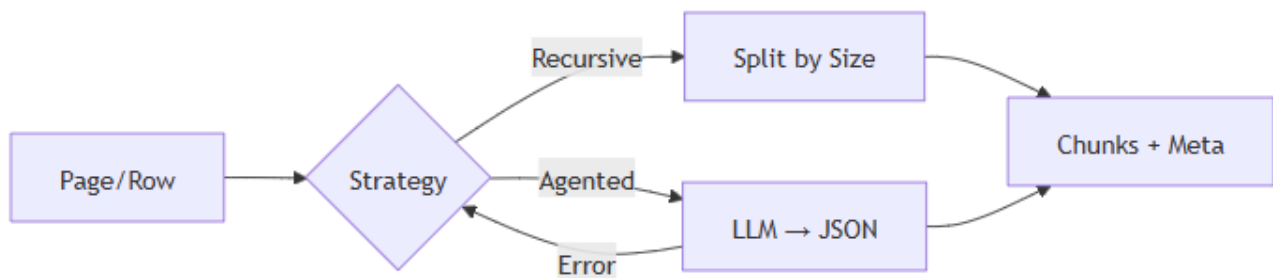


Figure 2: Chunking Strategies

- **Recursive Strategy (Default):** Reliable splitting using `RecursiveCharacterTextSplitter`. Tries to keep sentences/paragraphs intact within token limits (`max_tokens_per_chunk`, `chunk_overlap`).
- **Agented Strategy (Experimental):** Asks the LLM itself to identify logical breaks. Potentially smarter, but slower and requires careful error handling (falls back to Recursive).

3. Vectorizing Reality

To find relevant information quickly, we convert text chunks into numerical representations (embeddings) and store them in a searchable index.

3.1. Creating Embeddings (Embedder)

Transforms text into meaningful vectors using a pre-trained model.

- **Model:** `nomic-ai/nomic-embed-text-v1` via sentence-transformers. Chosen for strong retrieval performance and local usability.
- **Workflow:**
 1. Initialize the SentenceTransformer model.
 2. Input: List of text chunks from the Chunker.
 3. Process: Feed each chunk's text into `model.encode()`.
 4. Output: Add the resulting numerical vector (as a list) to the chunk's data structure ('embedding': [0.1, 0.2, ...]).




Figure 3: Embedding Process

3.2. Building the Vector Memory (VectorDB with ChromaDB)

A local, persistent database optimized for finding similar vectors (and thus, semantically similar text).

Setup:

- Uses `chromadb.PersistentClient` with `persist_dir` for local storage.
- Creates/loads a collection with `metadata={"hnsw:space": "cosine"}` for **cosine similarity** and fast HNSW indexing.

Storing Chunks:

1. **Input:** Text chunks + embeddings.
2. **Prep:** Assign unique `uuids` and ensure metadata has only simple types (string, number, bool).
3. **Store:** Use `collection.add()` with IDs, documents, embeddings, and metadata.

Retrieving Chunks:

1. **Input:** `query_text` + `n_results`.
2. **Embed:** Convert query to vector using the same Nomic model.
3. **Query:** Use `collection.query()` to find top similar vectors via cosine similarity.
4. **Output:** Returns list of matching chunks with text, metadata, and similarity score.




Figure 4: Vector Database Query Flow

4. Core Application: RAG Q&A (RAGPipeline)

This is where we answer questions by combining document knowledge with LLM intelligence.

4.1. The RAG Cycle: Retrieve -> Augment -> Generate

This three-step dance is fundamental to providing grounded answers.

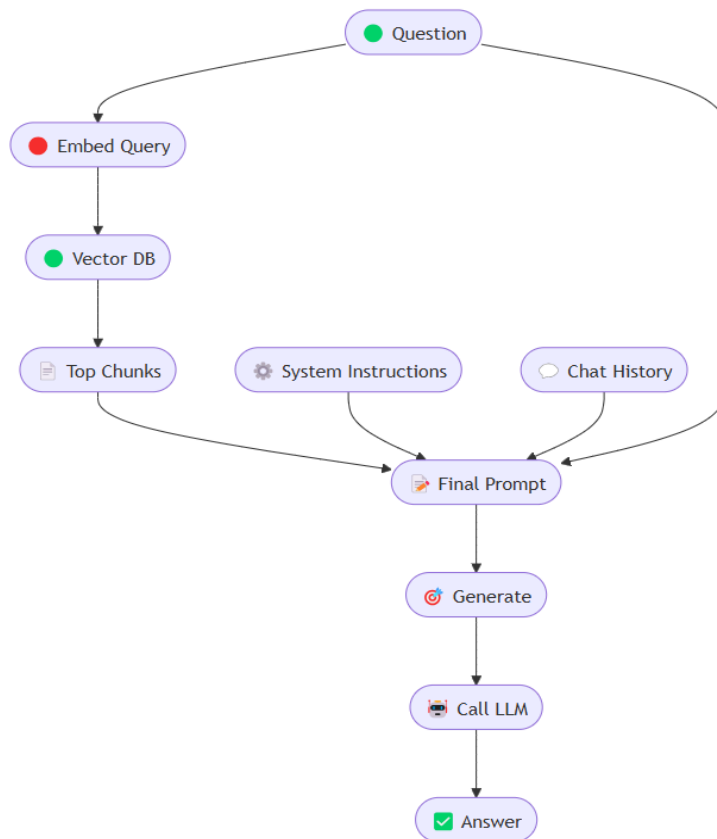


Figure 5: RAG Workflow

4.2. Algorithm Breakdown

1. **Retrieve:** Embed the user's question (Nomic model). Query VectorDB to fetch the `n_results` most relevant text chunks based on vector similarity.
2. **Augment:**
 - **Context is King:** Create a context string by combining the text (`page_content`) and key metadata (`source file`, `page/row`) from the retrieved chunks.
 - **Prompt Engineering:** Construct the final prompt for the LLM:
 - **System Role:** Use `SYSTEM_PROMPT_TEMPLATE` (from `config.py`) to tell the LLM its persona (e.g., "You are Dr. X's files assistant...") and inject the context string.
 - **Conversation:** Optionally, include recent `HumanMessage/AIMessage` pairs from `chat_history` to allow follow-up questions.
 - **The Ask:** Add the current user question as a `HumanMessage`.
3. **Generate:**
 - Send the assembled list of messages to the local LLM (`self.llm.invoke(messages)`).
 - The LLM uses the provided context and instructions to generate an informed answer.
 - Return the LLM's response text.

4.3. Conversational Context

If `use_history` is enabled in `config.py`, the system keeps track of the last few conversation turns (`chat_history`). This history is included in the "Augment" step, allowing the LLM to understand follow-up questions like "Tell me more about *that*."

5. Generative Tasks: Translation & Summarization

Beyond Q&A, we use the LLM's creative power for translation and generating concise summaries.

5.1. Translation (Translator)

Leverages the LLM to translate document content chunk by chunk.

- **Workflow:**

1. Load & Chunk the source document.
2. For each text chunk:
 - Create a specific prompt: "Translate the following text to {language}. Keep formatting... Text: {chunk_text}".
 - Invoke the LLM with this prompt.
- 3.
4. Assemble translated chunks into an output file.



Figure 6: Translation Workflow

- **Challenge:** Reliably preserving complex formatting depends heavily on the LLM's ability to follow instructions precisely.

5.2. Summarization (Summarizer)

Creates concise overviews using various strategies.

- **Core Strategies (via Langchain load_summarize_chain):**

1. **map_reduce:** Good for long docs. Summarize chunks -> Combine summaries.
2. **refine:** Good for flow. Iteratively update summary with each chunk.
3. **rerank (Custom):** Summarize small pieces -> Score/Rank -> Combine top N.
4. **Prompt Control:** Use different map/combine templates ('default', 'analytical', 'extractive') from config.py to change summary style.

- **Optional Clustering Workflow:**

1. **Group:** If enabled, use KMeans clustering on chunk embeddings to group similar content.
2. **Visualize:** Generate a t-SNE plot to show cluster separation.
3. **Summarize Clusters:** Apply the chosen strategy (MapReduce, etc.) to each cluster *individually*.
4. **Combine:** Create a final summary by summarizing the *cluster summaries*.




Figure 7: Summarization Workflow (with Optional Clustering)

6. Measuring Success: Evaluation & Performance

How good are the summaries? How fast is the system? We need metrics.

6.1. Summary Quality (SummaryEvaluator)

Uses the standard ROUGE metric to compare generated summaries against human references.

- **Algorithm:**
 1. Input: Generated summary, Reference summary (from summary_files_organic/).
 2. Tool: rouge_score.
 3. Calculate: ROUGE-1 (unigrams), ROUGE-2 (bigrams), ROUGE-L (longest common subsequence) - providing Precision, Recall, F1-Score for each.
 4. Log & Save: Record scores using Logger and save detailed results (CSV/JSON) in logs/evaluation_results/.

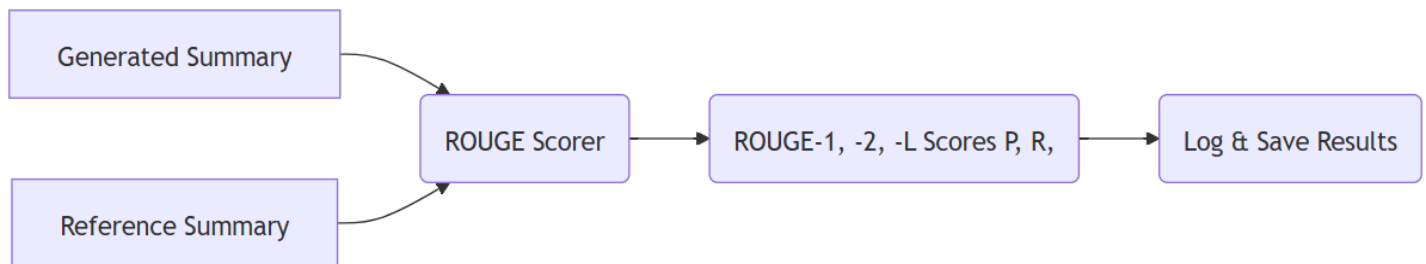


Figure 8: ROUGE Evaluation Flow

6.2. System Performance (Logger)

Tracks execution speed, focusing on LLM efficiency.

- **Algorithm:**
 1. **Time It:** Wrap key operations (especially llm.invoke, embedder.encode) with time.time() start/end calls.
 2. **Count Tokens:** For LLM outputs, use tiktoken to count the number of generated tokens.
 3. **Calculate:** Compute duration = end - start and tokens_per_second = output_tokens / duration.
 4. **Log It:** Use logger.metrics(operation, start, end, extra_info={'tokens/sec': ..., 'output_tokens': ...}) to record results.
- **Benefit:** Provides concrete data on how fast the local LLM performs specific tasks on the host machine.

6.3. Creative Touches

- **Agented Chunking:** Exploring LLM-native segmentation.
- **Smart Table Handling:** Translating tables to prose for better LLM comprehension.
- **Cluster-Driven Summaries:** Thematic grouping for structured summaries.
- **Flexible Summarization:** Multiple methods and prompt styles for tuning.

6.4. Final Thoughts

The Dr. X Analysis Suite is a robust, locally-focused NLP toolkit. It effectively tackles document diversity, implements advanced RAG and generative techniques, and provides mechanisms for evaluating both quality (ROUGE) and efficiency (Tokens/Sec). While constrained by local hardware and text-only analysis, it provides a powerful, transparent, and extensible platform for digging into Dr. X's textual legacy.